



Docker Deep Dive

A photograph of a person standing on a rocky cliff edge, looking out over a vast, calm ocean under a cloudy, overcast sky.

The Docker Engine

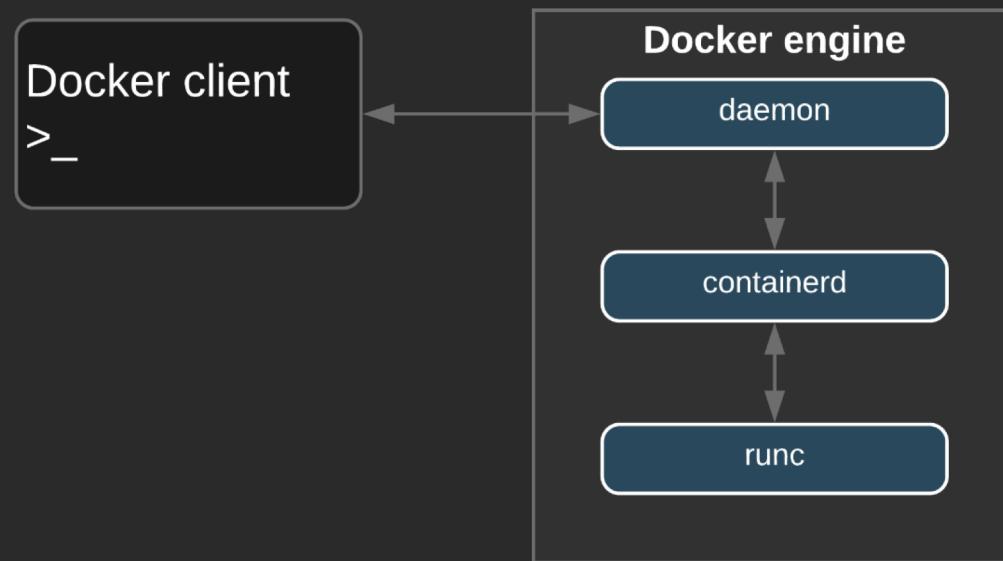
Under The Hood

Docker engine :

- Modular in design:
 - Batteries included but replaceable
- Based on an open-standards outline by the Open Container Initiative
- The major components:
 - Docker client
 - Docker daemon
 - `containerd`
 - `runc`
- The components work together to create and run containers



Under The Hood (cont.)



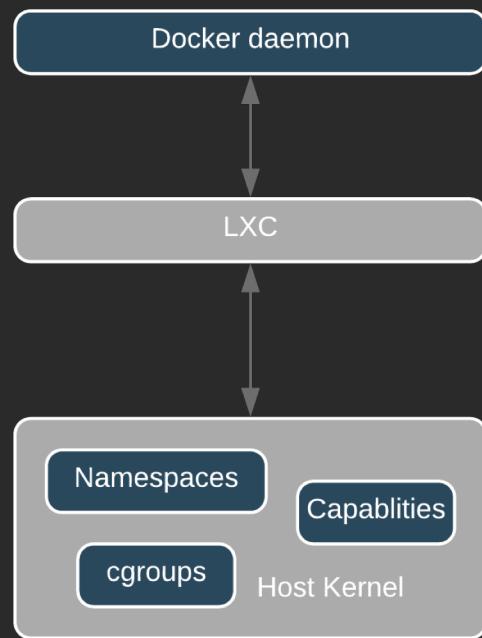
A Brief History of the Docker Engine

The first release of Docker

- The Docker daemon
 - Docker client
 - Docker daemon
 - Monolithic binary
 - Docker client
 - Docker API
 - Container runtime
 - Image builds
 - Much more...
- LXC
 - Namespaces
 - Control groups (cgroups)
 - Linux-specific



A Brief History of the Docker Engine (cont.)



Refactoring of the Docker Engine

LXC was later replaced with libcontainer:

- Docker 0.9
- Platform agnostic

Issues with the monolithic Docker daemon:

- Harder to innovate
- Slow
- Not what the ecosystem wanted

Docker became more modular:

- Smaller more specialized tools
- Pluggable architecture



Refactoring of the Docker Engine (cont.)

Open Container Initiative:

- Image spec
- Container runtime spec
- Version 1.0 released in 2017
- Docker, Inc. heavily contributed
- Docker 1.11 (2016) used the specification as much as possible



Refactoring of the Docker Engine (cont.)

runc

- Implementation of the OCI container-runtime-spec
- Lightweight CLI wrapper for libcontainer
- Create containers

containerd

- Manage container lifecycle:
 - Start
 - Stop
 - Pause
 - Delete
- Image management
- Part of the 1.11 release



Refactoring of the Docker Engine (cont.)

shim

- Implementation of daemonless Containers
- `containerd` forks an instance of `runc` for each new container
- `runc` process exits after the container is created
- `shim` process becomes the container parent
- Responsible for:
 - STDIN and STDOUT
 - Reporting exit status to the Docker daemon



Running Containers

```
docker container run -it --name <NAME> <IMAGE>:<TAG>
```

Creating a container:

- Use the CLI to execute a command.
- Docker client uses the appropriate API payload.
- The command POSTs to the correct API endpoint.
- The Docker daemon receives instructions.
- The Docker daemon calls `containerd` to start a new container.
- The Docker daemon uses gRPC, a CRUD style API.
- `containerd` creates an OCI bundle from the Docker image.
- Tells `runc` to create a container using the OCI bundle
- `runc` interfaces with the OS kernel to get the constructs needed to create a container



Running Containers (cont.)

Creating a container:

- This includes namespaces, `cgroups`, etc.
- The container process is started as a child process
- Once the container starts `runc` will exit
- Which completes the process and the container is now running



Process Summary

